Control Net for Image and Video Generation

[1] "Adding Conditional Control to Text-to-Image Diffusion Models" (Zhang et al, ICCV 2023)

[2] "Control-A-Video: Controllable Text-to-Video Diffusion Models with Motion Prior and Reward Feedback Learning" (Chen et al, 2023)

Overall Architecture

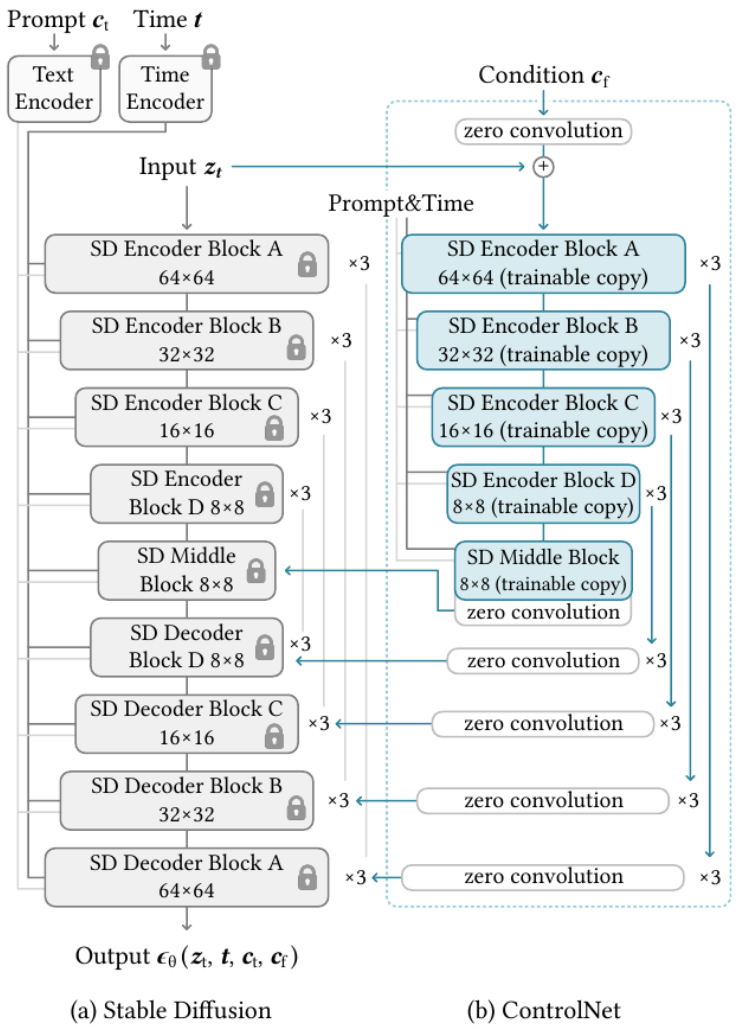

(a) Stable Diffusion (b) ControlNet

Sudden Convergence Phenomenon

Test input | training step 100 | step 1000 | step 2000

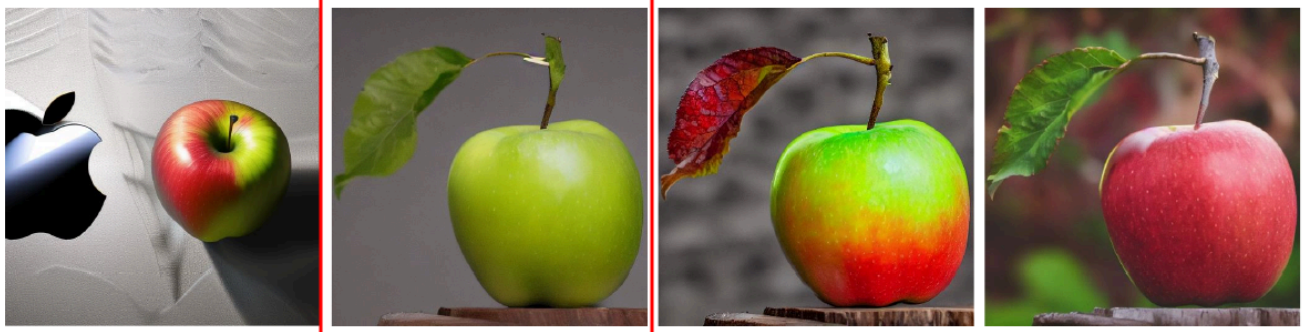

step 6100 | **step 6133** | step 8000 | step 12000

The model does **NOT gradually learn the control conditions** but abruptly succeeds in following the input conditioning image; usually in less than 10K optimization steps.

Classifier-Free Guidance

$$\epsilon_{\text{prd}} = \epsilon_{\text{uc}} + \beta_{\text{cfg}}(\epsilon_{\text{c}} - \epsilon_{\text{uc}})$$

where

- ϵ_{prd} — the model's final output
- β_{cfg} — the weight of guidance
- $\epsilon_{\text{uc}} = \texttt{UnetWithoutControlNet}(\mathbf{z}_t, t, \texttt{""}; \theta)$ — unconditional output
- $\epsilon_{\text{c}} = \texttt{UnetWithConrolNet}(\mathbf{z}_t, t, \mathbf{c}_t, \mathbf{c}_f; \theta, \phi)$ — conditional output
- $\mathbf{z}_t$ — noisy latent image
- $\mathbf{c}_t$ — text prompt
- $\mathbf{c}_f$ — condition image
- θ — pre-trained diffusion model's weights
- ϕ — control-net's weight

Classifier-Free Guidance Resolution Weighting

(a) Input Canny map (b) W/o CFG (c) W/o CFG-RW (d) Full (w/o prompt)

Classifier-Free Guidance Resolution Weighting

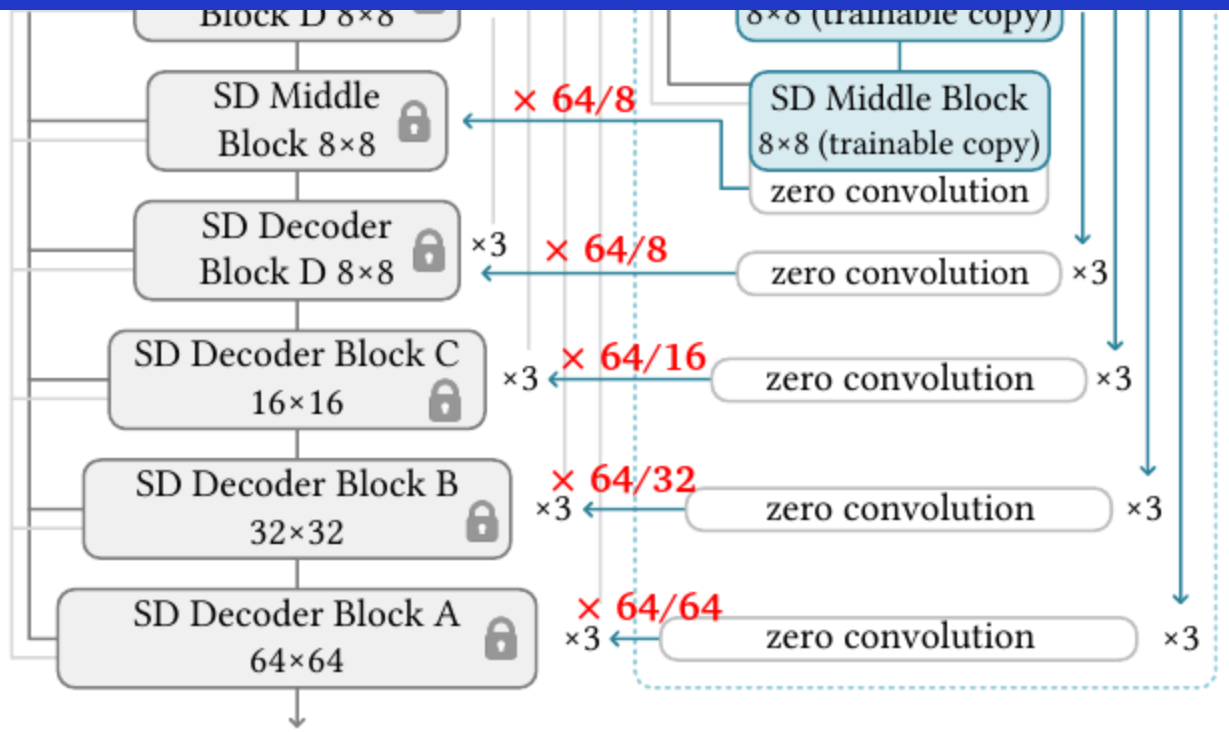

By multiplying a weight w_i to each connection between Stable Diffusion and ControlNet according to the resolution of each block $w_i = 64/h_i$, where h_i is the size of ith block,

e.g., $h_1 = 8, h_2 = 16, \ldots, h_{13} = 64$, we can achieve the better generation result.

ControlNet for VideoGeneration

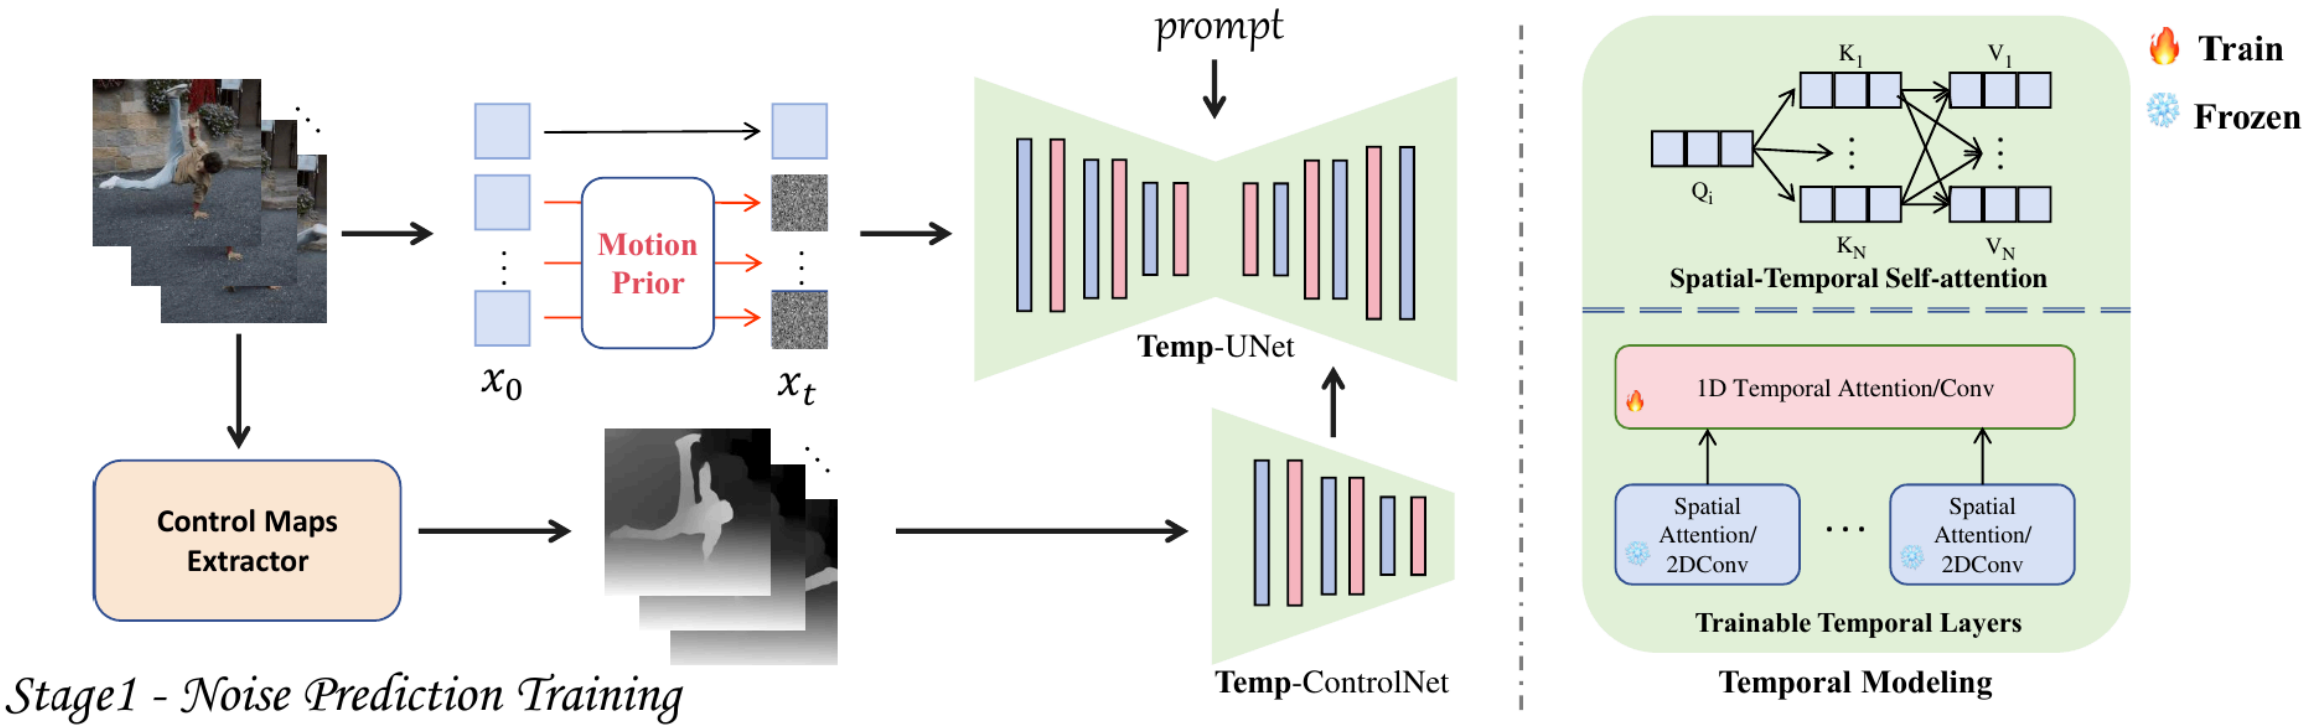

"Control-A-Video: Controllable Text-to-Video Diffusion Models with Motion Prior and Reward Feedback Learning" (Chen et al, 2023)